

# Introduction to Structured Prediction

He He

CDS, NYU

March 30, 2021

## Example: Part-of-speech (POS) Tagging

- Given a sentence, give a part of speech tag for each word:

$x$	$\underbrace{[\text{START}]}_{x_0}$	$\underbrace{\text{He}}_{x_1}$	$\underbrace{\text{eats}}_{x_2}$	$\underbrace{\text{apples}}_{x_3}$
$y$	$\underbrace{[\text{START}]}_{y_0}$	$\underbrace{\text{Pronoun}}_{y_1}$	$\underbrace{\text{Verb}}_{y_2}$	$\underbrace{\text{Noun}}_{y_3}$

- $\mathcal{V} = \{\text{all English words}\} \cup \{[\text{START}], ", ."]\}$
- $\mathcal{X} = \mathcal{V}^n, n = 1, 2, 3, \dots$  [Word sequences of any length]
- $\mathcal{P} = \{\text{START, Pronoun, Verb, Noun, Adjective}\}$
- $\mathcal{Y} = \mathcal{P}^n, n = 1, 2, 3, \dots$  [Part of speech sequence of any length]

# Multiclass Hypothesis Space

- **Discrete** output space:  $\mathcal{Y}(x)$   $\mathcal{Y} = \{1, \dots, k\}$ 
  - Very large but has structure, e.g., linear chain (sequence labeling), tree (parsing)
  - Size depends on input  $x$
- Base Hypothesis Space:  $\mathcal{H} = \{h: \mathcal{X} \times \mathcal{Y} \rightarrow \mathbb{R}\}$ 
  - $h(x, y)$  gives **compatibility score** between input  $x$  and output  $y$
- Multiclass hypothesis space

$$\mathcal{F} = \left\{ x \mapsto \arg \max_{y \in \mathcal{Y}} h(x, y) \mid h \in \mathcal{H} \right\}$$

- Final prediction function is an  $f \in \mathcal{F}$ .
- For each  $f \in \mathcal{F}$  there is an underlying compatibility score function  $h \in \mathcal{H}$ .

# Structured Prediction

- Part-of-speech tagging

$x$ :	he	eats	apples
$y$ :	pronoun	verb	noun

- Multiclass hypothesis space:

$$h(x, y) = w^T \Psi(x, y) \tag{1}$$

$$\mathcal{F} = \left\{ x \mapsto \arg \max_{y \in \mathcal{Y}} h(x, y) \mid h \in \mathcal{H} \right\} \tag{2}$$

- A special case of multiclass classification
- How to design the feature map  $\Psi$ ? What are the considerations?

# Unary features

- A **unary feature** only depends on
  - the label at a **single position**,  $y_i$ , and  $x$

- Example:

$$f(y_i, x)$$

$$\phi_1(x, y_i) = 1(x_i = \text{runs})1(y_i = \text{Verb})$$

$$\phi_2(x, y_i) = 1(x_i = \text{runs})1(y_i = \text{Noun})$$

$$\phi_3(x, y_i) = 1(x_{i-1} = \text{He})1(x_i = \text{runs})1(y_i = \text{Verb})$$

$$1(x_{i-2} = \text{<S>})$$

$$1(x_{i-1} \text{ has suffix } y)$$

# Markov features

- A **markov feature** only depends on
  - two **adjacent** labels,  $y_{i-1}$  and  $y_i$ , and  $x$
- Example:

$$\theta_1(x, y_{i-1}, y_i) = 1(y_{i-1} = \text{Pronoun})1(y_i = \text{Verb})$$

$$\theta_2(x, y_{i-1}, y_i) = 1(\underline{y_{i-1}} = \text{Pronoun})1(y_i = \text{Noun})$$

- Reminiscent of Markov models in the output space
- Possible to have higher-order features

# Local Feature Vector and Compatibility Score

- At each position  $i$  in sequence, define the **local feature vector** (unary and markov):

$$\Psi_i(x, y_{i-1}, y_i) = (\phi_1(x, y_i), \phi_2(x, y_i), \dots, \theta_1(x, y_{i-1}, y_i), \theta_2(x, y_{i-1}, y_i), \dots)$$

- And **local compatibility score** at position  $i$ :  $\langle w, \Psi_i(x, y_{i-1}, y_i) \rangle$ .
- The compatibility score for  $(x, y)$  is the sum of local compatibility scores:

$$\sum_i \langle w, \Psi_i(x, y_{i-1}, y_i) \rangle = \left\langle w, \sum_i \Psi_i(x, y_{i-1}, y_i) \right\rangle = \langle w, \Psi(x, y) \rangle, \quad (3)$$

where we define the **sequence feature vector** by

$$\Psi(x, y) = \sum_i \Psi_i(x, y_{i-1}, y_i). \quad \text{decomposable}$$

# Structured perceptron

Given a dataset  $\mathcal{D} = \{(x, y)\}$ ;

Initialize  $w \leftarrow 0$ ;

**for**  $iter = 1, 2, \dots, T$  **do**

**for**  $(x, y) \in \mathcal{D}$  **do**

$\hat{y} = \arg \max_{y' \in \mathcal{Y}(x)} w^T \psi(x, y')$ ;

**if**  $\hat{y} \neq y$  **then** // We've made a mistake

$w \leftarrow w + \Psi(x, y)$ ; // Move the scorer towards  $\psi(x, y)$

$w \leftarrow w - \Psi(x, \hat{y})$ ; // Move the scorer away from  $\psi(x, \hat{y})$

**end**

**end**

**end**

Identical to the multiclass perceptron algorithm except the  $\arg \max$  is now over the structured output space  $\mathcal{Y}(x)$ .



# Structured hinge loss

- Recall the generalized hinge loss

$$\ell_{\text{hinge}}(y, \hat{y}) \stackrel{\text{def}}{=} \max_{y' \in \mathcal{Y}(x)} (\Delta(y, y') + \langle w, (\Psi(x, y') - \Psi(x, y)) \rangle) \quad (4)$$

$= \begin{cases} 1 & y \neq y' \\ 0 & y = y' \end{cases}$

- What is  $\Delta(y, y')$  for two sequences?

- Hamming loss** is common:

$$\Delta(y, y') = \frac{1}{L} \sum_{i=1}^L 1(y_i \neq y'_i)$$

where  $L$  is the sequence length.

$\begin{matrix} 0 & 0 & 1 & 1 \\ 0 & 0 & 1 & 0 \end{matrix} \rightarrow \Delta(y, y') = 1$

- Can generalize to the cost-sensitive version using  $\delta(y_i, y'_i)$

## Exercise:

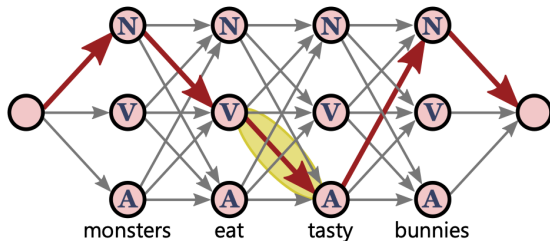
- Write down the objective of structured SVM using the structured hinge loss.
- Stochastic sub-gradient descent for structured SVM (similar to HW3 P3)
- Compare with the structured perceptron algorithm

# The argmax problem for sequences [BONUS]

**Problem** To compute predictions, we need to find  $\arg\max_{y \in \mathcal{Y}(x)} \langle w, \Psi(x, y) \rangle$ , and  $|\mathcal{Y}(x)|$  is exponentially large.

**Observation**  $\Psi(x, y)$  decomposes to  $\sum_i \Psi_i(x, y)$ .  $\rightarrow y_i, y_{i-1}$

**Solution** Dynamic programming (similar to the Viterbi algorithm)



What's the running time?

# The argmax problem in general

Efficient problem-specific algorithms:

problem	structure	algorithm
constituent parsing	binary trees with context-free features	CYK
dependency parsing	spanning trees with edge features	Chu-Liu-Edmonds
image segmentation	2d with adjacent-pixel features	graph cuts

General algorithm:

- Integer linear programming (ILP)

$$\max_z a^T z \quad \text{s.t. linear constraints on } z \quad (5)$$

- $z$ : indicator of substructures, e.g.,  $\mathbb{I}\{y_i = \text{article and } y_{i+1} = \text{noun}\}$
- constraints:  $z$  must correspond to a valid structure

## Multiclass algorithms

- Reduce to binary classification, e.g., OvA, AvA, ECCO
  - Good enough for simple multiclass problems
- Generalize binary classification algorithms using multiclass loss
  - Useful for problems with extremely large output space, e.g., structured prediction
  - Related problems: ranking, multi-label classification