

# Introduction to Structured Prediction

He He

CDS, NYU

March 30, 2021

2021-03-30

DS-GA 1003

Introduction to Structured Prediction

He He

CDS, NYU

March 30, 2021

# Example: Part-of-speech (POS) Tagging

- Given a sentence, give a part of speech tag for each word:

$x$	$\underbrace{[\text{START}]}_{x_0}$	$\underbrace{\text{He}}_{x_1}$	$\underbrace{\text{eats}}_{x_2}$	$\underbrace{\text{apples}}_{x_3}$
$y$	$\underbrace{[\text{START}]}_{y_0}$	$\underbrace{\text{Pronoun}}_{y_1}$	$\underbrace{\text{Verb}}_{y_2}$	$\underbrace{\text{Noun}}_{y_3}$

- $\mathcal{V} = \{\text{all English words}\} \cup \{[\text{START}], " . " \}$
- $\mathcal{X} = \mathcal{V}^n, n = 1, 2, 3, \dots$  [Word sequences of any length]
- $\mathcal{P} = \{\text{START, Pronoun, Verb, Noun, Adjective}\}$
- $\mathcal{Y} = \mathcal{P}^n, n = 1, 2, 3, \dots$  [Part of speech sequence of any length]

2021-03-30

## Example: Part-of-speech (POS) Tagging

Example: Part-of-speech (POS) Tagging

Given a sentence, give a part of speech tag for each word:

$x$	$\underbrace{[\text{START}]}_{x_0}$	$\underbrace{\text{He}}_{x_1}$	$\underbrace{\text{eats}}_{x_2}$	$\underbrace{\text{apples}}_{x_3}$
$y$	$\underbrace{[\text{START}]}_{y_0}$	$\underbrace{\text{Pronoun}}_{y_1}$	$\underbrace{\text{Verb}}_{y_2}$	$\underbrace{\text{Noun}}_{y_3}$

- $\mathcal{V} = \{\text{all English words}\} \cup \{[\text{START}], " . " \}$
- $\mathcal{X} = \mathcal{V}^n, n = 1, 2, 3, \dots$  [Word sequences of any length]
- $\mathcal{P} = \{\text{START, Pronoun, Verb, Noun, Adjective}\}$
- $\mathcal{Y} = \mathcal{P}^n, n = 1, 2, 3, \dots$  [Part of speech sequence of any length]

# Multiclass Hypothesis Space

- **Discrete** output space:  $\mathcal{Y}(x)$ 
  - Very large but has structure, e.g., linear chain (sequence labeling), tree (parsing)
  - Size depends on input  $x$
- Base Hypothesis Space:  $\mathcal{H} = \{h : \mathcal{X} \times \mathcal{Y} \rightarrow \mathbb{R}\}$ 
  - $h(x, y)$  gives **compatibility score** between input  $x$  and output  $y$
- Multiclass hypothesis space

$$\mathcal{F} = \left\{ x \mapsto \arg \max_{y \in \mathcal{Y}} h(x, y) \mid h \in \mathcal{H} \right\}$$

- Final prediction function is an  $f \in \mathcal{F}$ .
- For each  $f \in \mathcal{F}$  there is an underlying compatibility score function  $h \in \mathcal{H}$ .

2021-03-30

DS-GA 1003

## └ Multiclass Hypothesis Space

Multiclass Hypothesis Space

- **Discrete** output space:  $\mathcal{Y}(x)$ 
  - Very large but has structure, e.g., linear chain (sequence labeling), tree (parsing)
  - Size depends on input  $x$
- Base Hypothesis Space:  $\mathcal{H} = \{h : \mathcal{X} \times \mathcal{Y} \rightarrow \mathbb{R}\}$ 
  - $h(x, y)$  gives **compatibility score** between input  $x$  and output  $y$
- Multiclass hypothesis space
$$\mathcal{F} = \left\{ x \mapsto \arg \max_{y \in \mathcal{Y}} h(x, y) \mid h \in \mathcal{H} \right\}$$
- Final prediction function is an  $f \in \mathcal{F}$ .
- For each  $f \in \mathcal{F}$  there is an underlying compatibility score function  $h \in \mathcal{H}$ .

- Part-of-speech tagging

$x$ :    he            eats    apples  
 $y$ :    pronoun    verb    noun

- Multiclass hypothesis space:

$$h(x, y) = w^T \Psi(x, y) \quad (1)$$

$$\mathcal{F} = \left\{ x \mapsto \arg \max_{y \in \mathcal{Y}} h(x, y) \mid h \in \mathcal{H} \right\} \quad (2)$$

- A special case of multiclass classification
- How to design the feature map  $\Psi$ ? What are the considerations?

2021-03-30

## Structured Prediction

contextual dependence, efficient argmax

- Part-of-speech tagging

$x$ :    he            eats    apples  
 $y$ :    pronoun    verb    noun

- Multiclass hypothesis space:

$$h(x, y) = w^T \Psi(x, y) \quad (1)$$

$$\mathcal{F} = \left\{ x \mapsto \arg \max_{y \in \mathcal{Y}} h(x, y) \mid h \in \mathcal{H} \right\} \quad (2)$$

- A special case of multiclass classification

- How to design the feature map  $\Psi$ ? What are the considerations?

- A **unary feature** only depends on
  - the label at a **single position**,  $y_i$ , and  $x$
- Example:

$$\phi_1(x, y_i) = 1(x_i = \text{runs})1(y_i = \text{Verb})$$

$$\phi_2(x, y_i) = 1(x_i = \text{runs})1(y_i = \text{Noun})$$

$$\phi_3(x, y_i) = 1(x_{i-1} = \text{He})1(x_i = \text{runs})1(y_i = \text{Verb})$$

## Unary features

- A unary feature only depends on
  - the label at a **single position**,  $y_i$ , and  $x$

- Example:

$$\begin{aligned}\phi_1(x, y_i) &= 1(x_i = \text{runs})1(y_i = \text{Verb}) \\ \phi_2(x, y_i) &= 1(x_i = \text{runs})1(y_i = \text{Noun}) \\ \phi_3(x, y_i) &= 1(x_{i-1} = \text{He})1(x_i = \text{runs})1(y_i = \text{Verb})\end{aligned}$$

- A **markov feature** only depends on
  - two **adjacent** labels,  $y_{i-1}$  and  $y_i$ , and  $x$

- Example:

$$\theta_1(x, y_{i-1}, y_i) = 1(y_{i-1} = \text{Pronoun})1(y_i = \text{Verb})$$

$$\theta_2(x, y_{i-1}, y_i) = 1(y_{i-1} = \text{Pronoun})1(y_i = \text{Noun})$$

- Reminiscent of Markov models in the output space
- Possible to have higher-order features

- A **markov feature** only depends on
  - two **adjacent** labels,  $y_{i-1}$  and  $y_i$ , and  $x$

- Example:

$$\begin{aligned}\theta_1(x, y_{i-1}, y_i) &= 1(y_{i-1} = \text{Pronoun})1(y_i = \text{Verb}) \\ \theta_2(x, y_{i-1}, y_i) &= 1(y_{i-1} = \text{Pronoun})1(y_i = \text{Noun})\end{aligned}$$

- Reminiscent of Markov models in the output space

- Possible to have higher-order features

# Local Feature Vector and Compatibility Score

- At each position  $i$  in sequence, define the **local feature vector** (unary and markov):

$$\Psi_i(x, y_{i-1}, y_i) = (\phi_1(x, y_i), \phi_2(x, y_i), \dots, \theta_1(x, y_{i-1}, y_i), \theta_2(x, y_{i-1}, y_i), \dots)$$

- And **local compatibility score** at position  $i$ :  $\langle w, \Psi_i(x, y_{i-1}, y_i) \rangle$ .
- The compatibility score for  $(x, y)$  is the sum of local compatibility scores:

$$\sum_i \langle w, \Psi_i(x, y_{i-1}, y_i) \rangle = \left\langle w, \sum_i \Psi_i(x, y_{i-1}, y_i) \right\rangle = \langle w, \Psi(x, y) \rangle, \quad (3)$$

where we define the **sequence feature vector** by

$$\Psi(x, y) = \sum_i \Psi_i(x, y_{i-1}, y_i). \quad \text{decomposable}$$

2021-03-30

DS-GA 1003

## Local Feature Vector and Compatibility Score

Local Feature Vector and Compatibility Score

At each position  $i$  in sequence, define the **local feature vector** (unary and markov):

$$\Psi_i(x, y_{i-1}, y_i) = (\phi_1(x, y_i), \phi_2(x, y_i), \dots, \theta_1(x, y_{i-1}, y_i), \theta_2(x, y_{i-1}, y_i), \dots)$$

And **local compatibility score** at position  $i$ :  $\langle w, \Psi_i(x, y_{i-1}, y_i) \rangle$ .

The compatibility score for  $(x, y)$  is the sum of local compatibility scores:

$$\sum_i \langle w, \Psi_i(x, y_{i-1}, y_i) \rangle = \left\langle w, \sum_i \Psi_i(x, y_{i-1}, y_i) \right\rangle = \langle w, \Psi(x, y) \rangle, \quad (3)$$

where we define the **sequence feature vector** by

$$\Psi(x, y) = \sum_i \Psi_i(x, y_{i-1}, y_i). \quad \text{decomposable}$$

# Structured perceptron

Given a dataset  $\mathcal{D} = \{(x, y)\}$ ;

Initialize  $w \leftarrow 0$ ;

**for**  $iter = 1, 2, \dots, T$  **do**

**for**  $(x, y) \in \mathcal{D}$  **do**

$\hat{y} = \arg \max_{y' \in \mathcal{Y}(x)} w^T \psi(x, y')$ ;

**if**  $\hat{y} \neq y$  **then** // We've made a mistake

$w \leftarrow w + \Psi(x, y)$  ; // Move the scorer towards  $\psi(x, y)$

$w \leftarrow w - \Psi(x, \hat{y})$  ; // Move the scorer away from  $\psi(x, \hat{y})$

**end**

**end**

**end**

Identical to the multiclass perceptron algorithm except the  $\arg \max$  is now over the structured output space  $\mathcal{Y}(x)$ .

2021-03-30

DS-GA 1003

└ Structured perceptron

Structured perceptron

```
Given a dataset  $\mathcal{D} = \{(x, y)\}$ ;  
Initialize  $w \leftarrow 0$ ;  
for  $iter = 1, 2, \dots, T$  do  
    for  $(x, y) \in \mathcal{D}$  do  
         $\hat{y} = \arg \max_{y' \in \mathcal{Y}(x)} w^T \psi(x, y')$ ;  
        if  $\hat{y} \neq y$  then // We've made a mistake  
             $w \leftarrow w + \Psi(x, y)$  ; // Move the scorer towards  $\psi(x, y)$   
             $w \leftarrow w - \Psi(x, \hat{y})$  ; // Move the scorer away from  $\psi(x, \hat{y})$   
        end  
    end  
end
```

Identical to the multiclass perceptron algorithm except the  $\arg \max$  is now over the structured output space  $\mathcal{Y}(x)$ .



# Structured hinge loss

- Recall the generalized hinge loss

$$\ell_{\text{hinge}}(y, \hat{y}) \stackrel{\text{def}}{=} \max_{y' \in \mathcal{Y}(x)} (\Delta(y, y') + \langle w, (\Psi(x, y') - \Psi(x, y)) \rangle) \quad (4)$$

- What is  $\Delta(y, y')$  for two sequences?
- Hamming loss** is common:

$$\Delta(y, y') = \frac{1}{L} \sum_{i=1}^L 1(y_i \neq y'_i)$$

where  $L$  is the sequence length.

- Can generalize to the cost-sensitive version using  $\delta(y_i, y'_i)$

2021-03-30

DS-GA 1003

## Structured hinge loss

Structured hinge loss

- Recall the generalized hinge loss

$$\ell_{\text{hinge}}(y, \hat{y}) \stackrel{\text{def}}{=} \max_{y' \in \mathcal{Y}(x)} (\Delta(y, y') + \langle w, (\Psi(x, y') - \Psi(x, y)) \rangle) \quad (4)$$

- What is  $\Delta(y, y')$  for two sequences?

- Hamming loss is common:

$$\Delta(y, y') = \frac{1}{L} \sum_{i=1}^L 1(y_i \neq y'_i)$$

where  $L$  is the sequence length.

- Can generalize to the cost-sensitive version using  $\delta(y_i, y'_i)$

Exercise:


- Write down the objective of structured SVM using the structured hinge loss.
- Stochastic sub-gradient descent for structured SVM (similar to HW3 P3)
- Compare with the structured perceptron algorithm

Exercise:

- Write down the objective of structured SVM using the structured hinge loss.
- Stochastic sub-gradient descent for structured SVM (similar to HW3 P3)
- Compare with the structured perceptron algorithm

# The argmax problem for sequences

- Problem** To compute predictions, we need to find  $\text{argmax}_{y \in \mathcal{Y}(x)} \langle w, \Psi(x, y) \rangle$ , and  $|\mathcal{Y}(x)|$  is exponentially large.
- Observation**  $\Psi(x, y)$  decomposes to  $\sum_i \Psi_i(x, y)$ .
- Solution** Dynamic programming (similar to the Viterbi algorithm)



What's the running time?

Figure by Daumé III. A course in machine learning. Figure 17.1.

2021-03-30

## The argmax problem for sequences

The argmax problem for sequences

**Problem** To compute predictions, we need to find  $\text{argmax}_{y \in \mathcal{Y}(x)} \langle w, \Psi(x, y) \rangle$ , and  $|\mathcal{Y}(x)|$  is exponentially large.

**Observation**  $\Psi(x, y)$  decomposes to  $\sum_i \Psi_i(x, y)$ .

**Solution** Dynamic programming (similar to the Viterbi algorithm)

What's the running time?

Figure by Daumé III. A course in machine learning. Figure 17.1.

Let  $K = |\mathcal{Y}|$ , DP runtime  $O(K^2L)$ ,  $m$ th order Markov feature has runtime  $O(K^mL)$ , naive runtime  $O(K^L)$ .

# The argmax problem in general

Efficient problem-specific algorithms:

problem	structure	algorithm
constituent parsing	binary trees with context-free features	CYK
dependency parsing	spanning trees with edge features	Chu-Liu-Edmonds
image segmentation	2d with adjacent-pixel features	graph cuts

General algorithm:

- Integer linear programming (ILP)

$$\max_z a^T z \quad \text{s.t. linear constraints on } z$$

(5)

- $z$ : indicator of substructures, e.g.,  $\mathbb{I}\{y_i = \text{article and } y_{i+1} = \text{noun}\}$
- constraints:  $z$  must correspond to a valid structure

2021-03-30

## The argmax problem in general

The argmax problem in general

Efficient problem-specific algorithms:

problem	structure	algorithm
constituent parsing	binary trees with context-free features	CYK
dependency parsing	spanning trees with edge features	Chu-Liu-Edmonds
image segmentation	2d with adjacent-pixel features	graph cuts

General algorithm:

- Integer linear programming (ILP)

$$\max_z a^T z \quad \text{s.t. linear constraints on } z$$
- $z$ : indicator of substructures, e.g.,  $\mathbb{I}\{y_i = \text{article and } y_{i+1} = \text{noun}\}$
- constraints:  $z$  must correspond to a valid structure

(5)

- Reduce to binary classification, e.g., OvA, AvA, ECCO
  - Good enough for simple multiclass problems
- Generalize binary classification algorithms using multiclass loss
  - Useful for problems with extremely large output space, e.g., structured prediction
  - Related problems: ranking, multi-label classification

## Multiclass algorithms

- Reduce to binary classification, e.g., OvA, AvA, ECCO
  - Good enough for simple multiclass problems
- Generalize binary classification algorithms using multiclass loss
  - Useful for problems with extremely large output space, e.g., structured prediction
  - Related problems: ranking, multi-label classification