

Decision Trees

He He

CDS, NYU

April 6, 2021

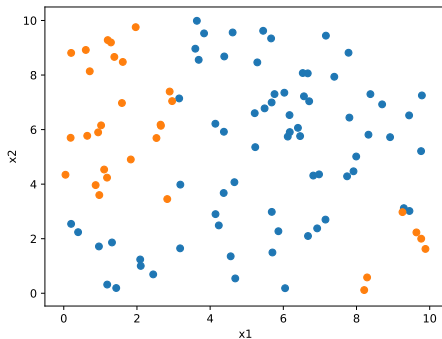
Today's lecture

- Our first inherently non-linear classifier: decision trees.
- Ensemble methods: bagging and boosting.

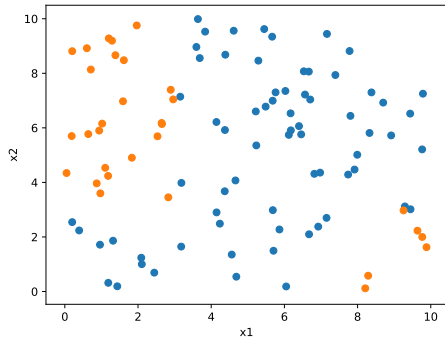
Decision Trees

Motivating example in 2d

- Partition data into different (axis-aligned) regions **recursively**



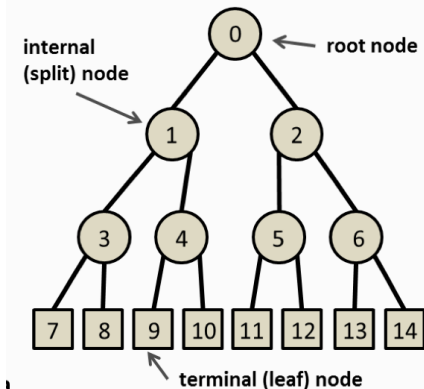
Classification flowchart



Is this a linear or non-linear classifier?

Decision trees setup

A general tree structure



We'll only consider

- *binary* trees (vs multiway trees where nodes can have more than 2 children)
- each node contains a subset of data points
- decisions at each node involve only a *single* feature (i.e. input coordinate)
- for continuous variables, splits always of the form

$$x_i \leq t$$

- for discrete variables, partitions values into two groups

Regularization of decision trees

- What will happen if we keep splitting the data?
 - Every data point will be in its own region—**overfitting**.
- When to stop splitting? (control complexity of the hypothesis space)
 - Limit number of total nodes.
 - Limit number of terminal nodes.
 - Limit tree depth.
 - Require minimum number of data points in a terminal node.

How to split

Goal Find a tree that minimize the task loss (e.g., squared loss) within a given complexity.

Problem Finding the optimal binary tree is computationally intractable.

Solution *Greedy* algorithm.

- Find the best split (according to some criteria) for a non-terminal node (initially the root)
- Add two children nodes
- Repeat until a stopping criterion is reached (e.g., max depth)

Evaluate splits

Let's think about what makes a good split.

Which one is better?

Split 1 $R_1 : 8+ / 2- \quad R_2 : 2+ / 8-$

Split 2 $R_1 : 6+ / 4- \quad R_2 : 1+ / 9-$

Which one is better?

Split 1 $R_1 : 8+ / 2- \quad R_2 : 2+ / 8-$

Split 2 $R_1 : 6+ / 4- \quad R_2 : 0+ / 10-$

In general, we want to produce *pure* nodes, i.e. close to single-class node.

Misclassification error in a node

Let's formalize things a bit.

- Consider classification case: $\mathcal{Y} = \{1, 2, \dots, K\}$.
- What's in a node?
 - Let node m represent region R_m , with N_m observations
 - Denote proportion of observations in R_m with class k by

$$\hat{p}_{mk} = \frac{1}{N_m} \sum_{\{i: x_i \in R_m\}} 1(y_i = k).$$

- Predict the majority class in node m :

$$k(m) = \arg \max_k \hat{p}_{mk}.$$

- Misclassification rate in node m :

$$1 - \hat{p}_{mk(m)}.$$

Node Impurity Measures

How to quantify impurity?

- Three measures of **node impurity** for leaf node m :

Misclassification error

$$1 - \hat{p}_{mk(m)}.$$

Gini index

$$\sum_{k=1}^K \hat{p}_{mk}(1 - \hat{p}_{mk})$$

Entropy / Information gain

$$-\sum_{k=1}^K \hat{p}_{mk} \log \hat{p}_{mk}.$$

- Gini index and entropy work well in practice.

Impurity of a split

A potential split produces two nodes, R_L and R_R . How do we score it?

- Suppose we have N_L points in R_L and N_R points in R_R .
- Let $Q(R_L)$ and $Q(R_R)$ be the node impurity measures for each node.
- Then find split that minimizes the *weighted average of node impurities*:

$$\frac{N_L Q(R_L) + N_R Q(R_R)}{N_L + N_R}$$

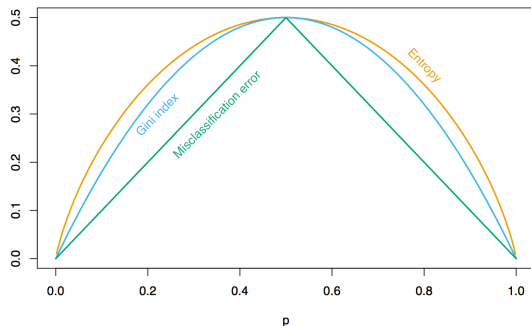
Example:

$R_1 : 8 + /2 -$ $R_2 : 1 + /4 -$

What's the weighted misclassification rate?

[discussion] Two-Class Node Impurity Measures

Consider binary classification. Let p be the relative frequency of class 1.



Misclassification error is not strictly concave thus may not guarantee improvement over the parent node.

Finding the Split Point

How to find a split point that minimizes a given impurity measure?

- Enumerate d features and $n - 1$ split points for each feature.
- Consider splitting on the j 'th feature x_j .
- If $x_{j(1)}, \dots, x_{j(n)}$ are the sorted values of the j 'th feature,
 - we only need to check split points between adjacent values
 - traditionally take split points halfway between adjacent values:

$$s_j \in \left\{ \frac{1}{2} (x_{j(r)} + x_{j(r+1)}) \mid r = 1, \dots, n-1 \right\}. \quad n-1 \text{ splits} \quad (1)$$

Regression trees

- Predict the mean value of a node

$$k(m) = \text{mean}(y_i \mid x_i \in R_m). \quad (2)$$

- Squared loss as the node impurity measure.
- Everything else remains the same as classification trees.

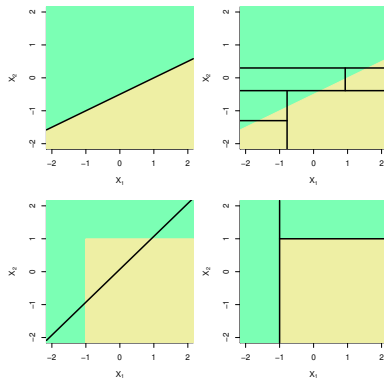
[discussion] Categorical features

- For a categorical feature, we split its values into two groups.
- Given a set of categories of size k , how many distinct splits? (its power set)
- Finding the optimal split is **intractable** in general.
- Approximations
 - Numeric encoding Randomly assign a number to each category
 - Binary classification: proportion of class 0
 - Regression: mean of targets of examples in the category, i.e. **mean encoding**
 - One-hot encoding May grow imbalanced trees, e.g., left-branching
 - Binary encoding Robust to large cardinality
- Statistical issues with categorical features
 - If a category has a very large number of categories, we can **overfit**.

- Trees are certainly easier to explain than other classifiers.
- Can be used to discover non-linear features.
- Small trees seem interpretable. For large trees, maybe not so easy.
- Approximate neural network decision boundaries to gain interpretability
 - Wu M, Hughes M, Parbhoo S, Zazzi M, Roth V, Doshi-Velez F. [Beyond Sparsity: Tree Regularization of Deep Models for Interpretability](#). Association for the Advancement of Artificial Intelligence (AAAI). 2018

Trees vs linear models

Trees have to work much harder to capture linear relations.



Review

Decision trees:

- *Non-linear* classifier that recursively partitions the input space.
- Non-metric: make no use of geometry, i.e. no inner-product or distances.
- Non-parametric: make no assumption of the data distribution.

Pros:

- Simple to understand.
- Interpretable, feature selection for free.

Cons:

- Poor linear modeling.
- Unstable / high variance, tend to **overfit**. → Next, how to fix this.